

Programming Assignment 1

Amrutanshu Mohanty 23B1208
Nenavath Sree Vamshi Madhav 23B1039
Dheeraj Kumar Maradana 23B0920
Arjoe Basak 23B1295

September 5, 2026

Datasets Overview

A brief summary of the dataset sizes evaluated in this assignment:

- **SciFact**: Corpus Size: 5,183 documents — Queries: 300
- **FEVER**: Corpus Size: 5,416,568 documents — Queries: 6,666
- **HotpotQA**: Corpus Size: 5,233,329 documents — Queries: 7,405
- **MSMARCO**: Corpus Size: 8,841,823 documents — Queries: 6908 (in dev split)

Part 1: Build an Index with Pyserini

Table 1 summarizes the build times and disk footprint of the Lucene inverted indexes for each dataset. The discrepancy for MSMARCO is because it was run on different machine

Dataset	Build Time (seconds)	Index Size
SciFact	12.60	10.09 MB
FEVER	850.36	4.39 GB
HotpotQA	423.77	2.27 GB
MSMARCO (Extra Credit)	372.65	4.02 GB

Table 1: Index build times and storage sizes using Pyserini.

Part 2: Sparse Retrieval Baselines

Table 2 reports the evaluation metrics for Default BM25, TF-IDF, and Tuned BM25.

Dataset	Method	nDCG@10	Recall@100	MRR@10	MAP	Latency (ms/q)
SciFact	Default BM25 ($k_1 = 0.9, b = 0.4$)	0.6789	0.9253	0.6457	0.6398	3.86 ms
	Classic TF-IDF	0.6744	0.9032	0.6390	0.6277	7.64 ms
	Tuned BM25 ($k_1 = 1.2, b = 0.75$)	0.6826	0.9282	0.6464	0.6396	1.74 ms
FEVER	Default BM25 ($k_1 = 0.9, b = 0.4$)	0.6513	0.9185	0.6230	0.5988	3.09 ms
	Classic TF-IDF	0.2075	0.6637	0.1727	0.1779	22.97 ms
	Tuned BM25 ($k_1 = 1.2, b = 0.1$)	0.6843	0.9275	0.6596	0.6316	4.27 ms
HotpotQA	Default BM25 ($k_1 = 0.9, b = 0.4$)	0.6330	0.7957	0.8032	0.5496	6.34 ms
	Classic TF-IDF	0.4376	0.6849	0.5668	0.3586	18.34 ms
	Tuned BM25 ($k_1 = 0.9, b = 0.4$)	0.6330	0.7957	0.8032	0.5496	2.90 ms

Table 2: Baseline sparse retrieval metrics and per-query retrieval latency.

Parameter Tuning Justification: The default BM25 implementation heavily penalizes very short or very long documents depending on the task. For *SciFact* and *HotpotQA*, reducing b limits the document-length normalization penalty, which benefits dense factual excerpts that may vary in length but contain highly concentrated relevant information. Lowering k_1 on *HotpotQA* further smooths out term-frequency saturation, allowing multiple distinct query terms to be matched more equitably. Conversely, *FEVER*’s Wikipedia passages are longer and more structured, benefiting a higher saturation threshold ($k_1 = 1.2$).

Experimental Setup: Dev Split Generation and Grid Search

To prevent data leakage, BM25 hyperparameter tuning was strictly performed on an independent development split before evaluating on the final held-out test splits. The dev splits were sourced as follows:

- **SciFact:** Because the BEIR distribution lacks a dedicated dev set for SciFact, the `beir/scifact/train` split was utilized for tuning.
- **FEVER and HotpotQA:** Sourced directly from their respective standard `beir/{dataset}/dev` splits, only randomly picked 1000 queries are used for finetuning
- **MSMARCO:** The standard `beir/msmarco/dev` split was randomly partitioned: 1,000 queries were sampled for parameter tuning, and the remaining queries were held out for the final test evaluation.

To ensure computational feasibility, any dev split exceeding 1,000 queries was randomly sub-sampled to exactly 1,000 queries using a fixed deterministic seed (seed=42).

Grid Search Methodology: We conducted an exhaustive grid search over a hyperparameter space defined by $k_1 \in \{0.3, 0.6, 0.9, 1.2, 1.6, 2.0\}$ and $b \in \{0.1, 0.2, 0.4, 0.6, 0.75, 0.9\}$, resulting in 36 unique configurations. For each configuration, the Lucene inverted index dynamically updated its BM25 similarity scoring formula, and a batch retrieval was executed over the dev queries. The retrieved rankings were scored via `pytrec_eval`, utilizing **nDCG@10** as the primary optimization objective. The (k_1, b) pair that maximized nDCG@10 on the dev set was decided upon and subsequently evaluated on the unseen test queries to produce the final reported metrics.

Part 3: Seeing Vocabulary Mismatch

Jaccard Overlap Analysis

We compared the Jaccard similarity (excluding stopwords) between queries and their gold-relevant documents. Table 3 displays the distributions for Successful queries (gold document in top-10) vs Failed queries (gold document not in top-10).

Dataset	Success Mean	Failure Mean	Success Median	Failure Median
SciFact	0.0514	0.0189	0.0460	0.0165
FEVER	0.0559	0.0330	0.0420	0.0222
HotpotQA	0.1438	0.0774	0.1250	0.0667

Table 3: Jaccard Overlap (Query vs Gold Documents) for Success vs. Failure.

Written Verdict: The empirical data overwhelmingly confirms that vocabulary mismatch is a primary driver of sparse retrieval failure. Across all three datasets, the mean and median Jaccard overlaps for successful queries are approximately 2 to 3 times higher than those for failed queries. Sparse retrievers depend completely on exact lexical matches; when the user’s phrasing diverges from the corpus phrasing, the BM25 score plummets.

Concrete Failure Case Examples and Modes




Figure 1: Jaccard Overlap Distributions for Successful vs. Failed queries at $k = 10$. Rows represent the text processing variations (Full set with stopwords, Full set without stopwords, and a 500-query Subset without stopwords). Columns represent the SciFact, FEVER, and HotpotQA datasets, respectively.

By analyzing the specific query-document pairs where BM25 failed to retrieve the gold document, we can categorize the vocabulary mismatch problem into several distinct failure modes:

Dataset: SciFact

1. **Query 619:** *"Increased vessel density along with a reduction in fibrosis decreases the efficacy of chemotherapy treatments."*
Gold Doc: "Hyaluronan impairs vascular function and drug delivery... stromal desmoplasia and vascular dysfunction..."
Failure Mode: Synonym / Terminology Mismatch. The query uses standard medical phrasing ("vessel density", "fibrosis", "chemotherapy"), but the document relies on parallel scientific terminology ("vascular function", "desmoplasia", "drug delivery"), resulting in a 0.0 Jaccard overlap.
2. **Query 870:** *"Obesity decreases life quality."*
Gold Doc: "Body-mass index and cause-specific mortality in 900,000 adults... associations of body-mass index (BMI)..."
Failure Mode: Concept Abstraction / Colloquial vs. Clinical Mismatch. BM25 cannot map the colloquial abstraction "Obesity" to the strict clinical measurement "body-mass index (BMI)" used in the document text.
3. **Query 1:** *"0-dimensional biomaterials show inductive properties."*
Gold Doc: "New opportunities: the use of nanotechnologies to manipulate and track stem cells..."
Failure Mode: Entity Granularity Mismatch. The query specifies the exact dimensional class ("0-dimensional biomaterials"), whereas the target document discusses the broader category ("nanotechnologies").
Lexical overlap is strictly zero.

Dataset: FEVER

1. **Query 168637:** *"The Thin Red Line (1998 film) portrays only frogs."*
Gold Doc: "Ben Chaplin (born Benedict John Greenwood ; 31 July 1969) is an English actor"
Failure Mode: Cross-Entity Evidence / Missing Context. FEVER claims are often verified by an indirect entity's page. The document proves Chaplin's involvement, but his biographical snippet doesn't explicitly repeat the movie title "The Thin Red Line" or the claim subject "frogs."
2. **Query 78985 (Evaluating against Gold Doc A):** *"Stomp the Yard stars a cat."*
Gold Doc A: "Brian Joseph White... is an American actor, producer, model..."
Failure Mode: Structural Knowledge Fragmentation. The BM25 algorithm relies on string matching, but the claim tests a factual relationship ("stars in"). The actor's introductory biography does not contain the literal string "Stomp the Yard," causing the rank to plummet.
3. **Query 78985 (Evaluating against Gold Doc B):** *"Stomp the Yard stars a cat."*
Gold Doc B: "Valarie Pettiford... is an American stage and television actress, dancer... received a Tony nomination..."
Failure Mode: Domain Specificity Dilution. Similar to Doc A, but further exacerbated by the introduction of competing domain terms ("stage", "television", "Tony nomination") which dilutes the TF-IDF weight of the document when searching for a specific movie title.

Dataset: HotpotQA

1. **Query 5a8c7595554299585d9e36b6:** "What government position was held by the woman who portrayed Corliss Archer in the film *Kiss and Tell*?"
Gold Doc: "Shirley Temple Black... American actress, singer... and diplomat..."
Failure Mode: Descriptive Reference vs Named Entity. The query describes an entity through a relation ("woman who portrayed Corliss Archer"). BM25 searches for "Corliss Archer", but the document only contains the answer entity ("Shirley Temple") and the target attribute ("diplomat").
2. **Query 5a879adb5542996e4f30887f:** "What other political position did the person who introduced the *DISCLOSE Act* hold?"
Gold Doc: "Christopher Van Hollen Jr... is the junior United States Senator from Maryland... U.S. Representative..."
Failure Mode: Multi-Hop Asymmetry. In multi-hop datasets, one document provides the bridging entity (who introduced the act) and the second provides the answer (Senator). BM25 fails on the second document because it lacks the original query context ("DISCLOSE Act").
3. **Query 5ae52c0f5542992663a4f12e:** "What major city is the Faith Lutheran Middle School and High School located by?"
Gold Doc: "Summerlin, Nevada... lies at the edge of the Spring Mountains and Red Rock Canyon... periphery of Las Vegas..."
Failure Mode: Missing Intermediate Entity. The target document verifies the location ("Las Vegas") of an intermediate entity ("Summerlin"), but contains zero mentions of the actual query subject ("Faith Lutheran Middle School"), making lexical retrieval impossible.

Part 4: Pseudo-Relevance Feedback

Part 4a: Rocchio and RM3 (Corpus-retrieved Feedback)

SciFact Metrics (Corpus PRF)

Method / Configuration	nDCG@10	Recall@100	MRR@10	MAP
Baseline Tuned BM25 (No PRF)	0.6826	0.9282	0.6464	0.6396
Pyserini Native RM3 ($N = 5, k = 10$)	0.6549	0.9103	0.6106	0.6067
Rocchio ($N = 3, k = 5$, pos-only)	0.5904	0.9030	0.5280	0.5225
Rocchio ($N = 5, k = 10$, pos-only)	0.5155	0.8773	0.4420	0.4380
Rocchio ($N = 10, k = 20$, pos-only)	0.4178	0.8967	0.3435	0.3437
Rocchio ($N = 5, k = 10$, pos+neg $\gamma = 0.15$)	0.5185	0.8807	0.4431	0.4390
Tuned Rocchio ($N = 3, k = 5, \beta = 0.25, \gamma = 0.0$)	0.5904	0.9030	0.5280	0.5225

Table 4: SciFact Pseudo-Relevance Feedback (Corpus PRF) metrics across configurations.

FEVER Metrics (Corpus PRF)

Method / Configuration	nDCG@10	Recall@100	MRR@10	MAP
Baseline Tuned BM25 (No PRF)	0.6843	0.9275	0.6596	0.6316
Pyserini Native RM3 ($N = 5, k = 10$)	0.6030	0.8985	0.5600	0.5352
Rocchio ($N = 3, k = 5$, pos-only)	0.5249	0.8641	0.4643	0.4466
Rocchio ($N = 5, k = 10$, pos-only)	0.4562	0.8475	0.3879	0.3726
Rocchio ($N = 10, k = 20$, pos-only)	0.2904	0.8109	0.2291	0.2277
Rocchio ($N = 5, k = 10$, pos+neg $\gamma = 0.15$)	0.4535	0.8490	0.3842	0.3698
Tuned Rocchio ($N = 3, k = 5, \beta = 0.25, \gamma = 0.0$)	0.5249	0.8641	0.4643	0.4466

Table 5: FEVER Pseudo-Relevance Feedback (Corpus PRF) metrics across configurations.

HotpotQA Metrics (Corpus PRF) HotpotQA Metrics (Corpus PRF)

Method / Configuration	nDCG@10	Recall@100	MRR@10	MAP
Baseline Tuned BM25 (No PRF)	0.6330	0.7957	0.8032	0.5496
Pyserini Native RM3 ($N = 5, k = 10$)	0.5457	0.7184	0.6960	0.4560
Rocchio ($N = 3, k = 5$, pos-only)	0.5340	0.7801	0.6260	0.4454
Rocchio ($N = 5, k = 10$, pos-only)	0.4733	0.7554	0.5342	0.3764
Rocchio ($N = 10, k = 20$, pos-only)	0.3609	0.7177	0.3979	0.2740
Rocchio ($N = 5, k = 10$, pos+neg $\gamma = 0.15$)	0.4761	0.7593	0.5356	0.3785
Tuned Rocchio ($N = 3, k = 5, \beta = 0.25, \gamma = 0.0$)	0.5340	0.7801	0.6260	0.4454

Table 6: HotpotQA Pseudo-Relevance Feedback (Corpus PRF) metrics across configurations.

Query Drift Qualitative Analysis:

Corpus PRF often suffers from severe query drift because the top-K retrieved documents contain confounding vocabulary that dilutes the original intent. **Query Drift Qualitative Analysis:** Corpus-based Pseudo-Relevance Feedback (PRF) often suffers from severe query drift because it blindly assumes the top-N retrieved documents are relevant. When the initial retrieval contains false positives, PRF extracts off-topic contextual terms that dilute the original query keywords, displacing the true gold documents down the ranking.

- **FEVER Drift Example (Query ID 177847):**

"Milk is based on the life of a person."

Baseline Rank: 1 $\rightarrow$ **Post-Rocchio Rank:** -1 (nDCG drops from 1.0 $\rightarrow$ 0.0)

Expansion terms: missing(0.20), company(0.12), granarolo(0.10), founded(0.08), cream(0.08), fresh(0.08)

Drift Analysis (Semantic Ambiguity): The user query refers to the biographical film *Milk*. However, the initial BM25 retrieval likely fetched documents about literal dairy milk and dairy corporations (e.g., "granarolo", "cream", "company"). Rocchio ingested these false-positive dairy terms and completely drifted the semantic focus away from the film, causing the gold document (*Milk*(film)) to fall entirely out of the top 100.

- **HotpotQA Drift Example (Query ID 5a861c0d5542994775f606f7):**

"C. Bernard Jackson founded the center that started the career of which actor, who played the part of Marty Castillo in 'Miami Vice'?"

Baseline Rank: 1 $\rightarrow$ **Post-Rocchio Rank:** 21 (nDCG drops from 1.0 $\rightarrow$ 0.0)

Expansion terms: delgado(0.13), gmrc(0.11), one(0.10), roger(0.09), play(0.07), dale(0.06), caesar(0.06)

Drift Analysis (Entity Hijacking): The query is highly specific and relies on intersecting multiple precise entities (C. Bernard Jackson, Miami Vice). The PRF mechanism extracted random proper nouns ("delgado", "roger", "caesar") from noisy biographical

documents in the top-K. These irrelevant entities hijacked the query weight, penalizing the documents that correctly matched the original actors and shows.

- **SciFact Drift Example (Query ID 660):**

"Ivermectin is used to treat onchocerciasis."

Baseline Rank: 1 $\rightarrow$ **Post-Rocchio Rank:** -1 (nDCG drops from 1.0 $\rightarrow$ 0.0)

Expansion terms: ascvd(0.07), arsb(0.06), chaseabc(0.04), spinal(0.03), light(0.03), cord(0.03)

Drift Analysis (Clinical Noise Injection): The baseline perfectly retrieved the document matching the two highly specific biomedical terms. However, Rocchio pulled in disjointed medical jargon ("ascvd", "spinal", "cord") from other clinical documents in the top-K. This completely diluted the focused intent of the query, resulting in a catastrophic drop in relevance.

Part 4b: HyDE (LLM-Generated Feedback)

SciFact Metrics (HyDE PRF vs Corpus PRF vs Naive)

Method / Configuration	nDCG@10	Recall@100	MRR@10	MAP
Baseline Tuned BM25	0.6826	0.9282	0.6464	0.6396
(A) Naive Concat HyDE ($N = 3, k = 3$)	0.6667	0.9516	0.6209	0.6148
(B) Rocchio-weighted HyDE ($N = 3, k = 3$)	0.6872	0.9337	0.6425	0.6344
(C) 4a Corpus Rocchio PRF ($N = 3, k = 3$)	0.5959	0.9070	0.5371	0.5323

FEVER Metrics (HyDE PRF vs Corpus PRF vs Naive)

Method / Configuration	nDCG@10	Recall@100	MRR@10	MAP
Baseline Tuned BM25	0.6843	0.9275	0.6596	0.6316
(A) Naive Concat HyDE ($N = 3, k = 3$)	0.6661	0.9199	0.6398	0.6102
(B) Rocchio-weighted HyDE ($N = 3, k = 3$)	0.6542	0.9151	0.6274	0.6011
(C) 4a Corpus Rocchio PRF ($N = 3, k = 3$)	0.5428	0.8868	0.4867	0.4678

HotpotQA Metrics (HyDE PRF vs Corpus PRF vs Naive)

Method / Configuration	nDCG@10	Recall@100	MRR@10	MAP
Baseline Tuned BM25	0.6330	0.7957	0.8032	0.5496
(A) Naive Concat HyDE ($N = 3, k = 3$)	0.5609	0.7803	0.6818	0.4826
(B) Rocchio-weighted HyDE ($N = 3, k = 3$)	0.6079	0.8034	0.7557	0.5260
(C) 4a Corpus Rocchio PRF ($N = 3, k = 3$)	0.5473	0.7874	0.6479	0.4597

Written Verdict on Feedback Source vs. Combination Method:

1. **Feedback Source Effect (HyDE vs Corpus PRF, Method = Rocchio):** Swapping the feedback source from retrieved corpus documents to LLM-generated hypothetical documents yields large, consistent improvements on every dataset tested: nDCG@10 rises by +0.09 on SciFact, +0.11 on FEVER, and +0.06 on HotpotQA at this setting, and the source effect stays positive across all 15 (N, k)-dataset combinations in our full sweep. Corpus PRF reinforces the errors of the first-stage retriever, since a bad initial ranking produces bad expansion terms. HyDE brings external world knowledge, generating relevant vocabulary that bridges the query-document gap independent of corpus bias.

2. Combination Method Effect (Rocchio vs Naive Concat, Source = HyDE):

Naive concatenation is competitive with Rocchio weighting, and the direction is dataset-dependent: Rocchio edges out naive concatenation on SciFact and HotpotQA at this conservative setting, but naive concatenation wins outright on FEVER. Across the full sweep the combination effect is positive in only 6 of 15 settings and turns negative on every dataset once k grows large, so whatever benefit Rocchio’s idf-weighted term selection provides is modest and conditional, not a reliable source of improvement on its own.

Conclusion: the feedback source is the dominant lever in this experiment. Its effect is large, positive, and consistent across every dataset and every configuration tested, while the combination method’s effect is smaller in magnitude, inconsistent in sign, and degrades as the expansion budget grows. Feedback source therefore matters substantially more than combination method for retrieval quality.//

Source Effect vs. Combination Effect (nDCG@10) Across Settings

Setting	Effect	SciFact	FEVER	HotpotQA
$N = 3, k = 3$	Source (B–C)	+0.0913	+0.1114	+0.0606
	Combination (B–A)	+0.0204	–0.0120	+0.0470
$N = 3, k = 5$	Source (B–C)	+0.0756	+0.1002	+0.0483
	Combination (B–A)	–0.0012	–0.0410	+0.0215
$N = 5, k = 10$	Source (B–C)	+0.1126	+0.0863	+0.0587
	Combination (B–A)	–0.0429	–0.1263	–0.0342

Impact of N and k : Increasing N , the number of hypothetical documents pooled per query, widens the feedback-source advantage, since pooling more sampled hallucinations gives Rocchio a steadier signal while corpus PRF weakens as more, noisier real documents enter its own feedback pool. Increasing k , the number of expansion terms kept, instead erodes the combination-method advantage: naive concatenation is invariant to k by construction, while Rocchio-weighted HyDE’s score falls as more terms are kept, turning the combination effect negative on every dataset by $k = 10$. We therefore adopt $N = 3, k = 3$ as our primary, most conservative setting, the point where the combination method’s benefit is best realized without undermining the far larger and more robust feedback-source effect.

Part 5: Sparse Retrieval with a Pretrained SPLADE Checkpoint

Metrics Comparison

SPLADE circumvents the two-stage overhead of traditional PRF by leveraging a Masked Language Modeling (MLM) head during the indexing phase to predict and assign weights to latent contextual terms.

Table 7 presents the performance of the `naver/splade-cocondenser-ensembledistil` checkpoint across the datasets, compared directly against our best Tuned BM25 baseline.

Dataset	Method	nDCG@10	Recall@100	MRR@10	MAP
SciFact	Tuned BM25 Baseline	0.6826	0.9282	0.6464	0.6396
	SPLADE (prebuilt index)	0.7036	0.9353	0.6723	0.6655
	Delta	+0.0210	+0.0071	+0.0259	+0.0259
FEVER	Tuned BM25 Baseline	0.6843	0.9275	0.6596	0.6316
	SPLADE (prebuilt index)	0.7879	0.9459	0.7864	0.7444
	Delta	+0.1036	+0.0184	+0.1268	+0.1128
HotpotQA	Tuned BM25 Baseline	0.6330	0.7957	0.8032	0.5496
	SPLADE (prebuilt index)	0.6868	0.8177	0.8655	0.6035
	Delta	+0.0538	+0.0220	+0.0623	+0.0539

Table 7: SPLADE Sparse Retrieval Metrics vs Baseline BM25.

Expansion Term Comparison (SPLADE vs Rocchio vs HyDE)

Table 8: Three-Way Expansion Term Comparison and Overlap (Scifact)

Query / Model	Top Generated Expansion Terms (and weights)
Query 1 (Scifact 1266): <i>"The risk of breast cancer among parous women increases with placental weight of pregnancies, and this association is strongest for premenopausal breast cancer."</i>	
SPLADE	pregnancy(1.53), tumor(1.39), nta(1.21), risks(1.07), pregnant(1.07), associations(1.05), increase(1.03), born(0.97), increased(0.92), female(0.80)
Rocchio	birth(0.03), bmi(0.03), american(0.03), white(0.02), incident(0.02), african(0.02), swedish(0.02), height(0.02)
HyDE	birth(0.05), placenta(0.04), given(0.04), exhibit(0.03), relationship(0.02), higher(0.02), wherein(0.02), heavier(0.02)
Overlaps	SPLADE $\cap$ Rocchio : 0.000 — SPLADE $\cap$ HyDE : 0.000 — Rocchio $\cap$ HyDE : 0.053
Query 2 (Scifact 1086): <i>"Sildenafil improves erectile function in men who experience sexual dysfunction as a result of the use of SSRI antidepressants."</i>	
SPLADE	erection(1.27), drug(1.26), sex(1.21), male(1.07), improve(0.97), treatment(0.95), depressed(0.91), orgasm(0.90), penis(0.86), medication(0.82)
Rocchio	natsal(0.03), health(0.03), women(0.02), serotonin(0.02), dhea(0.02), risk(0.02), sri(0.02), chlamydia(0.01)
HyDE	serotonin(0.11), cgmp(0.09), inhibitor(0.05), smooth(0.04), side(0.04), flow(0.04), cyclic(0.03), can(0.03)
Overlaps	SPLADE $\cap$ Rocchio : 0.000 — SPLADE $\cap$ HyDE : 0.000 — Rocchio $\cap$ HyDE : 0.053
Query 3 (Scifact 312): <i>"De novo assembly of sequence data has more specific contigs than unassembled sequence data."</i>	
SPLADE	assemble(1.18), mble(1.08), sequences(1.00), difference(0.90), sequencing(0.77), different(0.72), results(0.69), ructured(0.65), un(0.61), entry(0.59)
Rocchio	bias(0.07), rsem(0.05), aurea(0.04), seq(0.04), genlisea(0.03), length(0.03), full(0.02), rna(0.02)

Query / Model	Top Generated Expansion Terms (and weights)
HyDE	raw(0.10), dna(0.05), short(0.03), hand(0.03), longer(0.03), without(0.02), process(0.02), larger(0.02)
Overlaps	SPLADE $\cap$ Rocchio : 0.000 — SPLADE $\cap$ HyDE : 0.034 — Rocchio $\cap$ HyDE : 0.000
Query 4 (Scifact 239): <i>"Cellular aging closely links to an older appearance."</i>	
SPLADE	cell(1.79), age(1.49), old(1.37), link(1.23), appearances(1.10), related(1.03), elderly(0.99), appear(0.99), linked(0.73), earlier(0.69)
Rocchio	polyq(0.05), atr(0.04), thymic(0.03), will(0.03), cardiovascular(0.03), threshold(0.02), hypothesis(0.02), internet(0.02)
HyDE	skin(0.06), changes(0.04), gradual(0.03), process(0.02), stress(0.02), function(0.02), shorten(0.02), pro(0.01)
Overlaps	SPLADE $\cap$ Rocchio : 0.000 — SPLADE $\cap$ HyDE : 0.000 — Rocchio $\cap$ HyDE : 0.000
Query 5 (Scifact 213): <i>"CRP is not predictive of postoperative mortality following Coronary Artery Bypass Graft (CABG) surgery."</i>	
SPLADE	predicted(1.29), death(1.03), surgical(0.87), after(0.80), yes(0.79), transplant(0.77), cardiac(0.65), procedure(0.64), survival(0.62), heart(0.61)
Rocchio	tdm(0.06), statin(0.06), cardiac(0.06), aki(0.05), group(0.03), cpk(0.03), ima(0.03), dose(0.03)
HyDE	forecast(0.03), risk(0.03), marker(0.03), predictor(0.02), reflect(0.02), protein(0.02), reactant(0.02), post(0.01)
Overlaps	SPLADE $\cap$ Rocchio : 0.034 — SPLADE $\cap$ HyDE : 0.034 — Rocchio $\cap$ HyDE : 0.000
Query 6 (Scifact 1290): <i>"There is an inverse relationship between hip fractures and statin use."</i>	
SPLADE	fracture(1.40), correlation(1.09), yes(0.84), reverse(0.82), effect(0.72), drug(0.68), addiction(0.64), used(0.61), difference(0.59), broken(0.54)
Rocchio	bmd(0.06), bone(0.04), capsaicin(0.04), pain(0.03), risk(0.02), women(0.02), nonstatin(0.01), cgrp(0.01)
HyDE	cholesterol(0.07), hmg(0.06), bone(0.05), coa(0.05), research(0.03), lower(0.03), risk(0.03), class(0.03)
Overlaps	SPLADE $\cap$ Rocchio : 0.000 — SPLADE $\cap$ HyDE : 0.000 — Rocchio $\cap$ HyDE : 0.111
Query 7 (Scifact 124): <i>"Antiretroviral therapy reduces rates of tuberculosis across a broad range of CD4 strata."</i>	
SPLADE	tb(1.82), reduce(1.49), rate(1.22), hiv(1.21), treatment(1.03), reduction(1.00), drug(0.94), cds(0.89), reduced(0.76), vira(0.60)
Rocchio	art(0.09), hiv(0.08), partnership(0.04), count(0.03), haart(0.03), rifamycin(0.02), death(0.02), treatment(0.02)
HyDE	art(0.21), hiv(0.10), count(0.08), regardless(0.05), risk(0.03), crucial(0.02), shown(0.02), function(0.02)
Overlaps	SPLADE $\cap$ Rocchio : 0.071 — SPLADE $\cap$ HyDE : 0.071 — Rocchio $\cap$ HyDE : 0.176
Query 8 (Scifact 887): <i>"Only a minority of cells survive development after differentiation into stress-resistant spores."</i>	
SPLADE	stressed(1.39), cell(1.35), majority(1.22), resistance(1.20), pore(1.18), develop(1.08), survival(0.98), differentiate(0.71), few(0.67), pathogen(0.62)

Query / Model	Top Generated Expansion Terms (and weights)
Rocchio	pcna(0.07), foxo(0.05), spore(0.04), prestalk(0.03), discoideum(0.03), stalk(0.02), neutrophil(0.02), stem(0.02)
HyDE	spore(0.11), changes(0.05), process(0.05), dormant(0.04), proportion(0.03), cellular(0.03), cold(0.02), heat(0.02)
Overlaps	SPLADE $\cap$ Rocchio : 0.000 — SPLADE $\cap$ HyDE : 0.000 — Rocchio $\cap$ HyDE : 0.053
Query 9 (Scifact 1204): <i>"The combination of H3K4me3 and H3K79me2 is found in quiescent hair follicle stem cells."</i>	
SPLADE	hairs(1.40), cell(1.09), inhibitor(1.08), number(0.86), element(0.84), concentration(0.81), gene(0.81), enzyme(0.77), hormone(0.76), combinations(0.72)
Rocchio	aml(0.05), pcg(0.04), dermal(0.03), cryptic(0.03), self(0.02), transcription(0.02), nonskin(0.02), vivo(0.02)
HyDE	state(0.09), dormant(0.06), mark(0.05), chromatin(0.04), crucial(0.03), gene(0.02), phase(0.02), dna(0.02)
Overlaps	SPLADE $\cap$ Rocchio : 0.000 — SPLADE $\cap$ HyDE : 0.034 — Rocchio $\cap$ HyDE : 0.053
Query 10 (Scifact 628): <i>"Infection of human T-cell lymphotropic virus type 1 is most frequent in individuals of African origin."</i>	
SPLADE	africa(1.54), common(1.15), infections(1.10), pathogen(0.92), frequency(0.89), hiv(0.83), disease(0.81), viruses(0.76), cells(0.75), v(0.71)
Rocchio	hiv(0.10), htlv(0.08), leukaemia(0.03), retrovirus(0.03), viral(0.03), caribbean(0.03), sivagm(0.03), hsv(0.02)
HyDE	htlv(0.28), spastic(0.08), retrovirus(0.07), africa(0.06), lymphoma(0.05), west(0.05), leukemia(0.04), descent(0.04)
Overlaps	SPLADE $\cap$ Rocchio : 0.034 — SPLADE $\cap$ HyDE : 0.034 — Rocchio $\cap$ HyDE : 0.176

Table 9: Three-Way Expansion Term Comparison and Overlap (Fever)

Query / Model	Top Generated Expansion Terms (and weights)
Query 1 (Fever 26525): <i>"Shomu Mukherjee died at the age of 64."</i>	
SPLADE	death(1.31), die(1.29), old(1.18), india(1.08), jee(0.91), she(0.76), kumar(0.62), he(0.57), kolkata(0.56), khan(0.42)
Rocchio	samarth(0.34), tanuja(0.26), filmalaya(0.11), sashadhar(0.11), mukerji(0.09), produced(0.09), tanisha(0.08), kajol(0.08)
HyDE	information(0.23), claim(0.12), provide(0.11), official(0.10), historical(0.08), indian(0.08), academic(0.07), need(0.06)
Overlaps	SPLADE $\cap$ Rocchio : 0.000 — SPLADE $\cap$ HyDE : 0.034 — Rocchio $\cap$ HyDE : 0.000
Query 2 (Fever 48936): <i>"Reanimation is not a remix album."</i>	
SPLADE	albums(1.22), yes(1.03), parody(0.92), fake(0.84), remixes(0.81), copyright(0.78), song(0.71), music(0.65), genre(0.60), an(0.59)
Rocchio	released(0.28), linkin(0.11), hybrid(0.08), crunk(0.07), series(0.06), band(0.05), produced(0.05), mixed(0.05)

Query / Model	Top Generated Expansion Terms (and weights)
HyDE	electronic(0.35), original(0.11), house(0.09), project(0.09), series(0.08), music(0.08), producer(0.07), various(0.07)
Overlaps	SPLADE $\cap$ Rocchio : 0.000 — SPLADE $\cap$ HyDE : 0.071 — Rocchio $\cap$ HyDE : 0.053
Query 3 (Fever 10674): <i>"Manchester United F.C. was the world's highest-earning football club for a season."</i>	
SPLADE	salary(1.30), biggest(1.28), nfl(1.20), clubs(1.03), england(1.01), earn(0.94), fc(0.93), team(0.87), soccer(0.71), footballer(0.71)
Rocchio	city(0.16), one(0.11), cup(0.10), professional(0.07), considered(0.05), uefa(0.05), english(0.05), domestic(0.05)
HyDE	revenue(0.20), commercial(0.14), including(0.11), however(0.10), indeed(0.10), money(0.06), examine(0.06), claim(0.06)
Overlaps	SPLADE $\cap$ Rocchio : 0.034 — SPLADE $\cap$ HyDE : 0.000 — Rocchio $\cap$ HyDE : 0.000
Query 4 (Fever 3967): <i>"Yvonne Strahovski's real last name is Strzechowski."</i>	
SPLADE	she(1.03), poland(0.99), nationality(0.77), actor(0.75), founder(0.64), ki(0.52), lov(0.49), her(0.48), kowski(0.41), alias(0.39)
Rocchio	series(0.34), chuck(0.16), volkoff(0.08), miranda(0.06), however(0.06), mass(0.06), television(0.05), include(0.05)
HyDE	however(0.12), claim(0.06), adopted(0.06), professional(0.06), polish(0.05), munich(0.05), descent(0.05), common(0.05)
Overlaps	SPLADE $\cap$ Rocchio : 0.000 — SPLADE $\cap$ HyDE : 0.000 — Rocchio $\cap$ HyDE : 0.053
Query 5 (Fever 148057): <i>"The Nottingham Forest F.C. is barred from being a football club."</i>	
SPLADE	trent(1.61), forests(1.26), clubs(1.25), barr(1.19), nfl(1.14), england(0.83), ban(0.83), fc(0.74), association(0.74), sheffield(0.66)
Rocchio	history(0.31), formation(0.12), covers(0.12), information(0.11), competed(0.09), included(0.08), general(0.07), season(0.07)
HyDE	however(0.16), history(0.11), english(0.08), operate(0.08), placed(0.08), tier(0.07), professional(0.07), status(0.05)
Overlaps	SPLADE $\cap$ Rocchio : 0.000 — SPLADE $\cap$ HyDE : 0.000 — Rocchio $\cap$ HyDE : 0.111
Query 6 (Fever 76562): <i>"General Motors' non-automotive products divested in the 1980s through 2000s."</i>	
SPLADE	gm(1.40), vehicle(1.14), auto(1.11), motor(1.11), date(1.08), early(1.06), car(0.91), diving(0.90), company(0.84), closed(0.82)
Rocchio	company(0.18), engines(0.11), mini(0.10), cars(0.09), chevrolet(0.07), opel(0.07), ludvigsen(0.06), industrial(0.06)
HyDE	business(0.27), company(0.20), core(0.11), series(0.11), one(0.11), manufacturer(0.09), however(0.06), including(0.06)
Overlaps	SPLADE $\cap$ Rocchio : 0.034 — SPLADE $\cap$ HyDE : 0.071 — Rocchio $\cap$ HyDE : 0.111
Query 7 (Fever 60442): <i>"Tye Sheridan was cut from The Tree of Life."</i>	
SPLADE	trees(1.11), cutting(0.95), character(0.79), murder(0.77), film(0.69), died(0.66), actor(0.66), she(0.64), carved(0.54), death(0.43)
Rocchio	trees(0.24), one(0.06), represented(0.06), paper(0.05), forest(0.05), node(0.04), adaptation(0.03), experimental(0.03)

Query / Model	Top Generated Expansion Terms (and weights)
HyDE	malick(0.22), released(0.15), one(0.13), directed(0.13), film(0.11), jack(0.08), role(0.07), integral(0.06)
Overlaps	SPLADE $\cap$ Rocchio : 0.034 — SPLADE $\cap$ HyDE : 0.034 — Rocchio $\cap$ HyDE : 0.053
Query 8 (Fever 73630): <i>"The Selma to Montgomery marches were a failure to demonstrate the desire of African-American citizens to exercise their constitutional right to vote."</i>	
SPLADE	march(1.53), bam(1.37), voting(1.33), failed(1.10), marching(1.03), alabama(1.03), fail(0.97), rights(0.92), africa(0.91), because(0.85)
Rocchio	alabama(0.10), city(0.08), chestnut(0.07), sheyann(0.07), federal(0.05), one(0.05), march(0.04), several(0.04)
HyDE	march(0.10), series(0.07), lewis(0.07), television(0.06), civil(0.05), historical(0.05), outrage(0.05), alabama(0.05)
Overlaps	SPLADE $\cap$ Rocchio : 0.071 — SPLADE $\cap$ HyDE : 0.071 — Rocchio $\cap$ HyDE : 0.176
Query 9 (Fever 35754): <i>"Joseph Merrick was a man in the 19th century with perfect features."</i>	
SPLADE	joe(1.25), founder(1.05), early(0.70), celebrity(0.66), 1800s(0.65), actor(0.65), he(0.64), men(0.61), feature(0.54), his(0.54)
Rocchio	elephant(0.26), john(0.09), film(0.08), directed(0.07), company(0.06), received(0.06), adapted(0.06), hofsiss(0.06)
HyDE	body(0.27), appearance(0.25), physical(0.19), elephant(0.15), including(0.11), characterized(0.11), condition(0.10), carey(0.08)
Overlaps	SPLADE $\cap$ Rocchio : 0.000 — SPLADE $\cap$ HyDE : 0.000 — Rocchio $\cap$ HyDE : 0.053
Query 10 (Fever 207669): <i>"Andrew Carnegie has a 1889 article proclaiming "The Gospel of Wealth"."</i>	
SPLADE	rich(1.26), articles(1.14), church(0.99), early(0.91), fortune(0.91), proclaimed(0.85), founder(0.81), date(0.78), book(0.69), president(0.68)
Rocchio	company(0.21), comedy(0.07), wealthtv(0.07), one(0.07), directed(0.06), essay(0.06), television(0.06), material(0.06)
HyDE	distribution(0.15), essay(0.11), use(0.09), economic(0.07), moral(0.05), industrial(0.05), indeed(0.05), considered(0.05)
Overlaps	SPLADE $\cap$ Rocchio : 0.000 — SPLADE $\cap$ HyDE : 0.000 — Rocchio $\cap$ HyDE : 0.053

Table 10: Three-Way Expansion Term Comparison and Overlap (Hotpotqa)

Query / Model	Top Generated Expansion Terms (and weights)
Query 1 (Hotpotqa 5ae628fd5542995703ce8b33): <i>"Cinnamon sugar is a spice used to flavor this traditional Spanish and Portuguese fried pastry."</i>	
SPLADE	fry(1.38), spices(1.26), pie(1.22), glucose(1.18), spain(1.16), ingredient(0.90), food(0.77), baker(0.69), foods(0.66), taste(0.64)
Rocchio	khoshkar(0.16), reshteh(0.15), cruller(0.13), awameh(0.13), pie(0.12), dough(0.11), arabic(0.11), bean(0.11)

Query / Model	Top Generated Expansion Terms (and weights)
HyDE	aleo(0.16), sweet(0.11), preparation(0.10), ingredient(0.10), many(0.09), key(0.08), dough(0.08), dessert(0.07)
Overlaps	SPLADE $\cap$ Rocchio : 0.034 — SPLADE $\cap$ HyDE : 0.034 — Rocchio $\cap$ HyDE : 0.053
Query 2 (Hotpotqa 5a713ea95542994082a3e6e4): <i>"Alvaro Mexia had a diplomatic mission with which tribe of indigenous people?"</i>	
SPLADE	tribes(0.85), was(0.72), aboriginal(0.71), spain(0.67), native(0.66), diplomacy(0.63), ambassador(0.62), mexico(0.56), were(0.56), priest(0.53)
Rocchio	chumash(0.22), ynez(0.14), brazil(0.12), resident(0.12), santa(0.11), country(0.10), several(0.09), population(0.09)
HyDE	one(0.10), timucua(0.09), trade(0.08), taino(0.08), european(0.07), interaction(0.07), caribbean(0.06), cultural(0.06)
Overlaps	SPLADE $\cap$ Rocchio : 0.000 — SPLADE $\cap$ HyDE : 0.000 — Rocchio $\cap$ HyDE : 0.000
Query 3 (Hotpotqa 5ae022ab554299025d62a428): <i>"Which genus is native to arid or semi-desert areas, Alopecurus or Echeveria?"</i>	
SPLADE	species(1.20), genera(1.12), habitat(0.93), deserts(0.83), origin(0.41), plant(0.38), belong(0.38), subspecies(0.37), tree(0.33), breed(0.33)
Rocchio	species(0.33), strictiflora(0.33), setosa(0.18), mexican(0.15), plant(0.12), savior(0.11), mexico(0.11), okanagan(0.10)
HyDE	adapted(0.15), minimal(0.08), found(0.07), thrive(0.07), contrast(0.07), species(0.07), mexico(0.07), southwestern(0.06)
Overlaps	SPLADE $\cap$ Rocchio : 0.071 — SPLADE $\cap$ HyDE : 0.034 — Rocchio $\cap$ HyDE : 0.111
Query 4 (Hotpotqa 5ae28b76554299495565daa8): <i>"What three categories does income get levied under the Parliament of Australia's tax code?"</i>	
SPLADE	australian(1.37), category(1.18), taxes(1.11), types(0.98), congress(0.69), levy(0.60), accounting(0.47), codes(0.43), taxa(0.41), finance(0.41)
Rocchio	act(0.41), one(0.13), individual(0.08), revenue(0.08), business(0.07), important(0.07), personal(0.07), wealth(0.05)
HyDE	business(0.21), covers(0.12), company(0.07), professional(0.07), subject(0.05), exempt(0.05), purpose(0.05), certain(0.05)
Overlaps	SPLADE $\cap$ Rocchio : 0.000 — SPLADE $\cap$ HyDE : 0.000 — Rocchio $\cap$ HyDE : 0.053
Query 5 (Hotpotqa 5a7c9ead5542990527d554e4): <i>"As of 2010, what was the population of the county that was named after the 13th Speaker of the House of Representatives?"</i>	
SPLADE	karen(1.16), speakers(1.16), 13(1.03), congress(0.75), representative(0.64), senator(0.54), thirteenth(0.53), speech(0.52), vote(0.51), massachusetts(0.39)
Rocchio	census(0.15), blackford(0.14), macon(0.14), lick(0.12), one(0.12), maine(0.09), town(0.08), ohio(0.08)
HyDE	census(0.25), cannon(0.23), illinois(0.10), specific(0.09), rayburn(0.09), joseph(0.09), bureau(0.08), data(0.07)
Overlaps	SPLADE $\cap$ Rocchio : 0.000 — SPLADE $\cap$ HyDE : 0.000 — Rocchio $\cap$ HyDE : 0.053
Query 6 (Hotpotqa 5add21555542992ae4cec496): <i>"What is the start date of this tour that was the third concert tour by an American singer-songwriter and was completed in support of the album Red?"</i>	

Query / Model	Top Generated Expansion Terms (and weights)
SPLADE	three(1.08), dates(1.04), song(0.68), festival(0.58), singers(0.53), albums(0.53), supporting(0.47), tours(0.46), reddish(0.44), touring(0.42)
Rocchio	dear(0.10), swift(0.09), studio(0.08), truth(0.08), began(0.07), kany(0.07), boleto(0.06), williams(0.06)
HyDE	swift(0.29), featured(0.16), taylor(0.16), europe(0.12), released(0.09), included(0.08), america(0.06), stage(0.05)
Overlaps	SPLADE $\cap$ Rocchio : 0.000 — SPLADE $\cap$ HyDE : 0.034 — Rocchio $\cap$ HyDE : 0.053
Query 7 (Hotpotqa 5ab259b25542993be8fa98de): <i>"What word or phrase is found in both the history of Belgium and cockfighting?"</i>	
SPLADE	belgian(1.39), words(1.03), phrases(0.99), meaning(0.82), historical(0.81), fighting(0.80), fight(0.76), find(0.51), said(0.50), weapon(0.46)
Rocchio	many(0.17), defined(0.17), belgicism(0.11), adverb(0.11), useful(0.08), judicial(0.08), one(0.08), french(0.07)
HyDE	belgian(0.11), symbol(0.09), cockerel(0.08), cultural(0.08), worn(0.07), context(0.07), pigeon(0.06), many(0.05)
Overlaps	SPLADE $\cap$ Rocchio : 0.000 — SPLADE $\cap$ HyDE : 0.071 — Rocchio $\cap$ HyDE : 0.053
Query 8 (Hotpotqa 5a8aed1555429970aeb70334): <i>"What vehicle, made by UROVESA, is most like the four-wheel drive military light truck produced by AM General?"</i>	
SPLADE	model(0.98), vehicles(0.89), trucks(0.88), manufactured(0.73), army(0.59), manufacturer(0.58), similar(0.56), company(0.55), car(0.50), cars(0.42)
Rocchio	uro(0.20), jeep(0.12), vamtac(0.08), terrain(0.08), use(0.08), gun(0.08), purpose(0.07), specific(0.07)
HyDE	hmmwv(0.15), making(0.14), various(0.11), shares(0.09), eroc(0.08), several(0.08), robust(0.06), road(0.06)
Overlaps	SPLADE $\cap$ Rocchio : 0.034 — SPLADE $\cap$ HyDE : 0.000 — Rocchio $\cap$ HyDE : 0.053
Query 9 (Hotpotqa 5a85e24a5542994775f6067f): <i>"Are Hoodoo Gurus and Pierre Bowvier of the same nationality?"</i>	
SPLADE	different(0.71), ethnic(0.63), citizenship(0.50), he(0.43), character(0.42), religion(0.37), founder(0.32), origin(0.31), gay(0.31), yes(0.31)
Rocchio	released(0.41), produced(0.18), album(0.14), release(0.12), group(0.10), leilani(0.09), crank(0.08), rock(0.08)
HyDE	australian(0.14), formed(0.13), french(0.12), band(0.09), individual(0.09), rock(0.09), musical(0.08), information(0.08)
Overlaps	SPLADE $\cap$ Rocchio : 0.000 — SPLADE $\cap$ HyDE : 0.000 — Rocchio $\cap$ HyDE : 0.111
Query 10 (Hotpotqa 5a7c68575542996dd594b91c): <i>"What was the sequel of the game that e was published by U.S. Gold in 1992?"</i>	
SPLADE	sequels(1.31), us(1.30), 1991(1.28), games(1.08), publishing(1.03), film(0.84), early(0.81), 1993(0.76), date(0.68), successor(0.64)
Rocchio	developed(0.36), released(0.18), flashback(0.16), europe(0.15), strider(0.14), company(0.10), video(0.09), identity(0.09)
HyDE	question(0.22), released(0.20), provide(0.19), specific(0.15), information(0.12), answer(0.11), crimson(0.09), one(0.09)

Query / Model	Top Generated Expansion Terms (and weights)
Overlaps	SPLADE $\cap$ Rocchio : 0.034 — SPLADE $\cap$ HyDE : 0.034 — Rocchio $\cap$ HyDE : 0.053

Disagreement Analysis:

- **SciFact 1086:** SPLADE flawlessly identifies the latent semantic entities of the query (*Sildenafil* $\rightarrow$ *erection, penis, male*; *antidepressants* $\rightarrow$ *depressed*). Because SPLADE’s MLM head was pretrained on vast textual data, it understands the biological linkage independent of the target corpus. In contrast, Rocchio blindly grabs disconnected medical terms (*natsal, chlamydia, dhea*) that simply happened to co-occur in the highest-scoring corpus passages, leading to a strict Jaccard overlap of 0.000.
- **FEVER 26525:** For the query regarding Shomu Mukherjee’s death, Rocchio immediately drifts into tangentially related Bollywood actors and film studios (*kajol, tanuja, filmalaya*). SPLADE correctly abstracts the core event to general biographical facts (*death, die, old, kolkata*). HyDE, functioning as a generative LLM, produces self-evaluating phrasing (*information, claim, provide, official*), showing how generative models often hallucinate verification language rather than extracting strict indexing vocabulary.
- **FEVER 48936:** For the query about remix albums, Rocchio overfits to specific artist metadata present in the top-K corpus documents (*linkin, hybrid, crunk*). SPLADE ignores the metadata and explicitly abstracts the negative phrasing (“is not a remix”) into categorical semantic labels (*fake, parody, copyright*), demonstrating deep zero-shot understanding of the query intent.

Extra Credit 1: Document-Side Expansion (doc2query)

An alternative to expanding queries at retrieval time (query-side expansion) is to expand documents at indexing time (document-side expansion). Following the *doc2query* methodology, we utilized the `qwen2.5:7b-instruct` LLM to predict plausible pseudo-queries for each of the 5,183 documents in the SciFact corpus. These predictions were appended to the original document text, a new Lucene inverted index was compiled via Pyserini, and the Tuned BM25 baseline was re-evaluated on the expanded corpus.

Metrics and Trade-Off Analysis

Table 11 highlights the retrieval performance and computational resource trade-offs between the original unexpanded index and the doc2query expanded index.

Method (SciFact)	nDCG@10	Recall@100	Index Size	Query Latency
Original BM25	0.6826	0.9282	10.09 MB	6.07 ms
doc2query BM25	0.6917	0.9259	11.54 MB	4.54 ms

Table 11: Performance and resource trade-offs for Document-Side Expansion.

Resource Trade-off Discussion: Document-side expansion successfully improved the overall nDCG@10 from 0.6826 to 0.6917. Crucially, because all LLM generation is precomputed during the offline indexing phase, the system maintains ultra-low sub-millisecond retrieval latency at runtime (dropping slightly to 4.54 ms per query) with zero live LLM API calls. This

completely bypasses the severe runtime latency bottlenecks seen in query-side expansion techniques like HyDE.

The cost of this precomputation is paid in offline indexing overhead. Predicting 3 queries per document for the 5,183 SciFact documents required substantial GPU inference time upfront. Additionally, injecting these generated terms into the documents can swell up the final inverted index size on disk, although it didn't do so here (from 10.09 MB to 11.54 MB, a 1.14x space overhead). The index compilation step itself (via Pyserini) remained exceptionally fast, taking only 3.43 seconds once the corpus was expanded.

Part 3 Failure Case Recovery

We evaluated whether injecting LLM-predicted vocabulary into the documents at indexing time could bridge the hard vocabulary mismatches identified in Part 3 (where the original BM25 failed to retrieve the gold document in the Top 10 entirely).

- **Total Baseline Failed Queries (k=10):** 49
- **Recovered into Top-10 by doc2query:** 2 queries (4.08%)

While doc2query successfully improved the rank of already-retrieved documents (hence the overall nDCG boost), it was only able to completely recover 2 out of the 49 hard failure cases. For example, for Query ID 715, the LLM successfully generated pseudo-queries for the target document containing precise domain terminology ("*miRNA target abundance*," "*ceRNA hypothesis*"), enriching the sparse postings enough for BM25 to boost it from Rank 13 into Rank 4.

However, because the LLM is blindly predicting queries from the document text without seeing the actual user's query, it often fails to hallucinate the exact colloquialisms or specific abstractions a user might search with, which explains why query-side expansion (HyDE) remains slightly more targeted for vocabulary mismatch.

Extra Credit 2: Full-Scale MSMARCO Evaluation

To demonstrate pipeline scalability at massive web scale, we repeated Parts 1, 2, and 5 over the full *MSMARCO* passage corpus (8,841,823 documents).

Part 1 Indexing Characteristics

- **Corpus Size:** 8,841,823 passages
- **Evaluation Split:** 5,980 held-out queries (partitioned from the 6,980 `beir/msmarco/dev` queries; 1,000 queries were held out for (k_1, b) grid tuning)
- **Index Build Time:** 372.65 seconds (6.21 minutes via multithreaded Pyserini Lucene indexer)
- **On-Disk Index Size:** 4.02 GB (stored with positions, docvectors, and raw JSON payloads)

Part 2 Baseline Retrieval & Part 5 SPLADE Metrics

Table 12 summarizes the retrieval effectiveness and per-query latency across Default BM25, Classic TF-IDF, Tuned BM25 ($k_1 = 0.6, b = 0.75$), and the prebuilt SPLADE++ impact index (`msmarco-v1-passage.splade-pp-ed`).

Method / Configuration	nDCG@10	Recall@100	MRR@10	MAP	Latency (ms/q)
Default BM25 ($k_1 = 0.9, b = 0.4$)	0.2261	0.6606	0.1817	0.1897	3.99 ms
Classic TF-IDF	0.1720	0.6034	0.1350	0.1435	20.23 ms
Tuned BM25 ($k_1 = 0.6, b = 0.75$)	0.2348	0.6724	0.1877	0.1953	2.67 ms
SPLADE++ (Prebuilt Impact Index)	0.4466	0.9081	0.3804	0.3854	18.50 ms

Table 12: MSMARCO Sparse Retrieval Baselines (Part 2) vs SPLADE++ (Part 5) on 5,980 held-out dev queries.

Analysis: Hyperparameter tuning on MSMARCO favors lower $k_1 = 0.6$ (smoothing term saturation) and higher $b = 0.75$ (stronger length normalization for variable passage lengths), delivering improvements in both nDCG@10 (+0.0087) and search latency (2.67 ms/query). SPLADE++ dramatically outperforms all lexical baselines (+0.2118 nDCG@10 and +0.2357 Recall@100), demonstrating that learned neural expansions successfully overcome massive vocabulary mismatches in natural web search queries.

Extra Credit 3: Custom SPLADE Training and Domain Adaptation

Instead of relying solely on the off-the-shelf checkpoint from Part 5, we implemented and trained our own SPLADE models using an InfoNCE contrastive ranking loss with in-batch negatives combined with scheduled FLOPS sparsity regularization on both query and document vectors.

Evolution of the Base Encoder: Raw DistilBERT vs. Warm-Started SPLADE

1. **Initial Approach (Raw DistilBERT):** Following the original SPLADE formulation we initially trained from scratch starting with a standard masked language model **distilbert-base-uncased**.
 - On **SciFact** (5,183 documents, 1,109 queries), 5 epochs allowed the model to iterate over the entire corpus repeatedly, yielding a solid $nDCG@10 = 0.4837$ and $Recall@100 = 0.8294$.
 - On **FEVER** (5.4 million documents), training raw DistilBERT from scratch for 5 epochs on a 10,000-query subsample encountered severe cold-start representation collapse ($nDCG@10 = 0.1118$). Because the model only observed $< 0.1\%$ of the corpus without cross-encoder teacher distillation (MarginMSE), the FLOPS regularization penalty aggressively squashed token activations before contrastive ranking could stabilize, producing sparse but noisy term activations (Mean Jaccard Overlap with pretrained = 0.1360).
2. **Transition to Warm-Started SPLADE Base:** To scale domain adaptation to multi-million document corpora, we initialized the encoder from **naver/splade-cocondenser-ensembledistil** and fine-tuned it on FEVER claim-evidence pairs. Starting from mature retrieval weights preserved the structural sparsity and token gating discipline while allowing the model to adapt its expansion vocabulary to FEVER’s fact-verification domain without catastrophic rank degradation.

Retrieval Performance Comparison

Table 13 compares our custom fine-tuned models against the Part 5 Pretrained baseline.

Dataset	Model Variant	Base Initialization	nDCG@10	Recall@100	MRR@10	MAP
SciFact	Part 5 Pretrained Baseline	splade-cocondenser	0.7036	0.9353	0.6723	0.6655
	Custom Trained (5 Epochs)	distilbert-base	0.4837	0.8294	0.4409	0.4408
FEVER	Part 5 Pretrained Baseline	splade-cocondenser	0.7879	0.9459	0.7864	0.7444
	Custom Trained (5 Epochs)	distilbert-base	0.1118	0.1400	0.1110	0.1048
	Custom Fine-Tuned (3 Epochs)	splade-cocondenser	0.1291	0.1400	0.1306	0.1246

Table 13: Retrieval performance of Custom Trained SPLADE models vs Pretrained Part 5 baselines.

Expansion Term Agreement and Domain Specialization

We compared the top-10 expansion terms generated by the trained models against the pretrained Part 5 checkpoint across representative queries:

- **SciFact Biomedical Specialization (Query ID 3):**

"1,000 genomes project enables mapping of genetic sequence variation consisting of rare variants..."

Pretrained: genome(2.03), pen(1.43), projects(1.27), map(1.12), variations(1.11)

Fine-Tuned (SciFact): gene(1.44), dna(1.34), map(1.17), rna(1.04), transcript(0.56), variant(0.53)

Analysis: Training on SciFact successfully suppressed generic multi-meaning words (e.g., "projects") and shifted expansion mass toward dense biomedical indicators ("dna", "rna", "transcript"), demonstrating domain-specific term alignment.

- **FEVER Entity Expansion Stability (Query ID 70041):**

"2 Hearts is a musical composition by Minogue."

Pretrained: heart(2.18), kylie(1.65), two(1.64), artist(1.18), song(1.15), composed(0.94), music(0.89)

DistilBERT-Trained: heart(1.86), kylie(1.59), love(1.47), song(1.23), video(1.12), maxi(0.89), 2nd(0.84)

SPLADE-Fine-Tuned: heart(1.51), two(1.50), music(0.82), song(0.79), artist(0.79), composer(0.60), kylie(0.60)

Analysis: While the raw DistilBERT model drifted toward weak contextual tokens ("love", "video", "maxi"), the warm-started SPLADE model maintained a high Jaccard overlap of **0.667** with the pretrained model, refining musicology terms ("composer", "composed", "artist").

- **FEVER Political Entity Alignment (Query ID 163803):**

"Ukrainian Soviet Socialist Republic was a founding participant of the UN."

Pretrained: ukraine(1.99), nations(1.18), founded(1.16), russia(1.01), socialism(0.99), communism(0.89)

SPLADE-Fine-Tuned: ukraine(1.77), russia(1.00), nations(0.91), communism(0.86), russian(0.75), yugoslavia(0.56)

Analysis: Fine-tuning achieved a Jaccard overlap of **0.538**, discovering relevant historical geopolitical entities ("yugoslavia", "europe") essential for verifying 20th-century historical claims. Across 10 evaluation queries, switching from DistilBERT to the SPLADE base encoder increased mean expansion overlap from **0.1360** to **0.4496**.

LLM Chat Links

The chat histories can be viewed in the `chat_histories` folder inside our GitHub repository. The GitHub repository can be found [here](#).